

# Resilient Key Value Store Design

## 1. Introduction

The document describes the design of a replication based resilient key value store. Strong consistency and strict ordering of update operations on the different replicas are ensured using a transactional memory mechanism.

## 2. Application Interface

Users are provided with a HashMap like APIs for storing and manipulating their data. One place creates the map, and passes it to other places that wish to use it as follows:

```
val mapName = "my_map";
val map = DataStore.getInstance().makeResilientMap(mapName);

for ( p in Place.places() ) at (p) async {
    map.put ("some key", "some value");
}
```

A user can also perform multiple operations atomically within a single transaction as follows:

```
for ( p in Place.places() ) at (p) async {
    val txId = map.startTransaction();
    val x = map.get(txId, "X");
    val y = map.get(txId, "Y");
    map.put (txId, "Z", x+y);
    map.commit(txId);
}
```

The commit operation might fail when places perform conflicting transactions. In that case, the application need to repeat the transaction<sup>1</sup>.

## 3. Data Partitioning

The key/value records are partitioned among places. Currently, the number of partitions equals to `Place.numPlaces()`.<sup>2</sup> Consistent hashing is used for determining the specific key partition.

## 4. Topology-Aware Replication

Each partition is replicated in at least two places.

A [PartitionTable](#) at each place stores the Replicas (i.e. places) that store each partition.

A partition is not allowed to be replicated on places that exist in the same node.

Partition Id	The index of the below three arrays represents the partition id								
	0	1	2	3	4	5	6	7	8
Replica-1	Place(0)	Place(1)	...	...	...	...	...	Place(7)	Place(8)
Replica-2	Place(1)	Place(2)	...	...	...	...	...	Place(8)	Place(0)
Replica-3	Place(2)	Place(3)	...	...	...	...	...	Place(0)	Place(1)

## 5. Topology-Aware Leaders Assignment

A leader and a deputy leader places are selected to handle the loss of replicas.

Using the topology information, the leader and deputy leader places are selected from different nodes.

---

<sup>1</sup> Possible enhancement: perform automatic repetition upon failure by logging the issued map actions.

<sup>2</sup> Possible enhancement: allow configuring some places as Spare.

The following figure shows the main design modules.

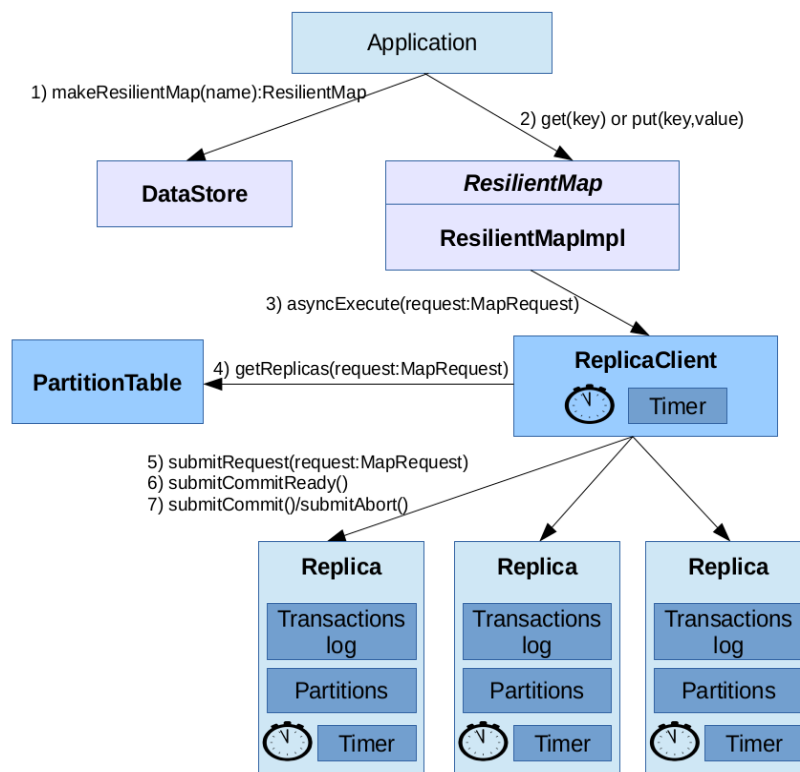


Fig. 1: Main design modules

## 6. Replica and ReplicaClient

Accessing the replicated key/value records is done using client-server protocols.

The **Replica** module, is responsible for the server side operations as it is the container of the actual data.

The **ReplicaClient** is the client side; it performs the map operations (e.g. get and put) on behalf of the application. Using the **PartitionTable** it identifies the replicas that can serve the application request and starts communicating with them using specific protocols that aim to ensure strong replicas consistency.

**Asynchronous communication:** the communication between a **Replica** and **ReplicaClient** is asynchronous. However, responses are expected within a specific time period, otherwise, the involved transaction is aborted.

At the **Replica** side, a late response from a client results in aborting the transaction.

At the **ReplicaClient** side, a late response from one of the replicas results in aborting the transaction.

**Dead Replica:** A **ReplicaClient** checks for the status of **Replicas** at the following times:

- Before starting a transaction, and
- Periodically, while waiting for a response from a replica

When it detects that one replica is dead, it notifies the leader place, terminates the transaction, and tries it again after some time.

**Dead ReplicaClient:** Each **Replica** periodically checks the status of **ReplicaClients** that has non-completed transactions. When a replica client is found dead, its transactions are aborted. The leader place is *not* notified.

## **7. Distributed Transactional Memory**

- versioned values
- isolation (store the updates in a local copy)
- distributed commitment protocol
  - when a transaction is ready to commit
  - abort my self vs abort others